

Kdump -very Briefly!

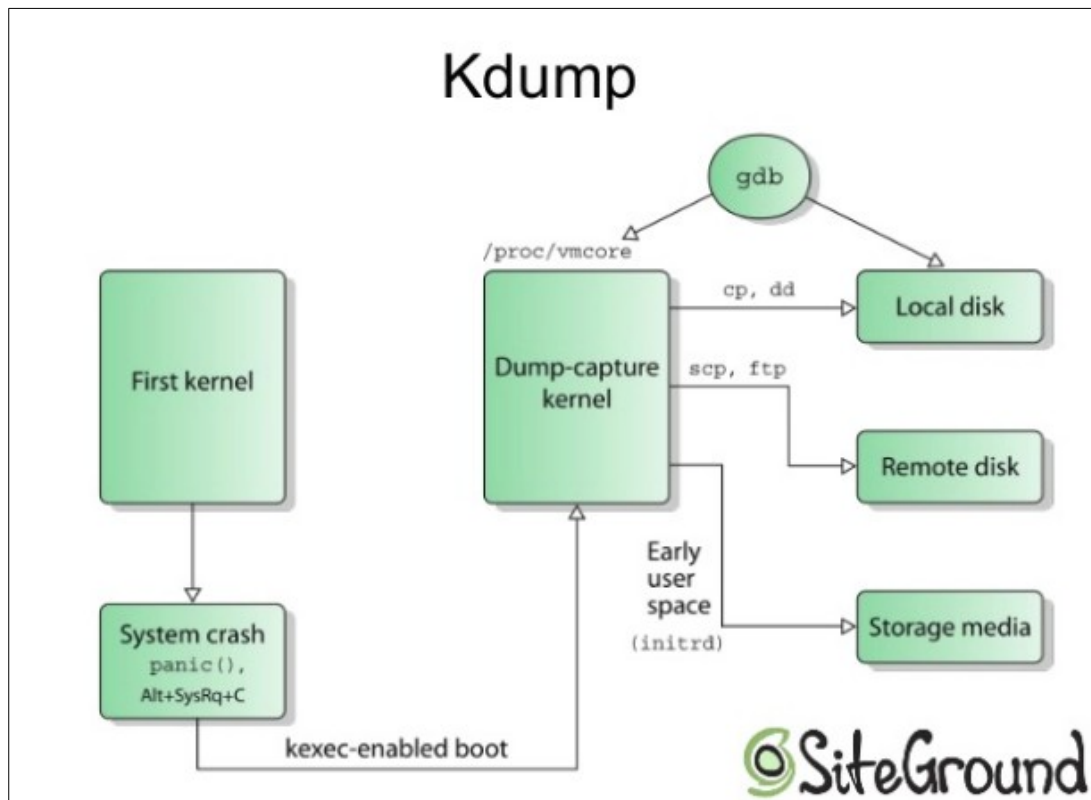
[Kernel Documentation for Kdump - The kexec-based Crash Dumping Solution](#)

[kdump \(Linux\)](#) on Wikipedia

Source: <https://www.slideshare.net/azilian/linux-kernel-crashdump>

- No dependencies, theoretically ideal, but...
 - Based on kexec
 - Not all arch support kexec
 - Not easy to setup
 - Boots a second kernel to retrieve the crash vmcore
 - Almost useless in cases of HW failure
 - Needs assistance of other tools for analysis





Tip: Analyze the kdump image with crash (instead of GDB)

1. Bootloader kernel command-line to first kernel:

```
console=ttyMXC0 rootfstype=ext4 root=/dev/mmcblk0 rw rootwait init=/sbin/init
crashkernel=128M@0x78000000
```

2. Just after first kernel has booted:

```
$ dmesg |grep -i crash
[ 0.000000] Reserving 128MB of memory at 1920MB for crashkernel (System RAM: 1784MB)
[ 0.000000] Kernel command line: console=ttyMXC0 rootfstype=ext4 root=/dev/mmcblk0
rw rootwait init=/sbin/init crashkernel=128M@0x78000000
$
```

Once *kexec* has successfully run on the original (first) kernel, can verify via `sysfs`

(`CONFIG_SYSFS` required for exactly this):

3. Before *kexec*

```
ARM / $ ls /sys/kernel/kexec_*
/sys/kernel/kexec_crash_loaded /sys/kernel/kexec_loaded
/sys/kernel/kexec_crash_size
ARM / $ cat /sys/kernel/kexec_*
0
134217728
0
ARM / $
```

4. After successful *kexec*

```
ARM / $ cat /sys/kernel/kexec_*
cat /sys/kernel/kexec_*
1
134217728
```

0
ARM / \$

5. Cause a panic!

echo c > /proc/sysrq-trigger

Demo on a Qemu-emulated Freescale i.MX6 platform ([more details here on the kaiwanTECH blog](#)). The first kernel crashes, and then, the dump-capture kernel boots!

```
ARM / $ id
uid=0 gid=0
ARM / $ echo c > /proc/sysrq-trigger    << or an actual kernel Oops/panic >>
[ 460.417261] sysrq: SysRq : Trigger a crash
[ 460.423293]
[ 460.424273] ===== << from lockdep, ignore for now >>
[ 460.424965] [ INFO: suspicious RCU usage. ]
[ 460.426708] 4.1.46 #2 Not tainted
[ 460.427276] -----
[ 460.427864] include/linux/rcupdate.h:570 Illegal context switch in RCU read-side critical
section!
[ 460.429056]
[ 460.429056] other info that might help us debug this:
[ 460.429056]
[ 460.430726]
[ 460.430726] rcu_scheduler_active = 1, debug_locks = 0
[ 460.432040] 3 locks held by sh/130:
[ 460.432717] #0: (sb_writers#4){.+.+.}, at: [<800fb398>] vfs_write+0x140/0x15c
[ 460.437331] #1: (rcu_read_lock){.....}, at: [<8031c664>] __handle_sysrq+0x0/0x254
[ 460.438777] #2: (&mm->mmap_sem){+++++}, at: [<8001ecb0>] do_page_fault+0x78/0x37c
[ 460.440515]
[ 460.440515] stack backtrace:    << back to the crash >>
[ 460.441491] CPU: 0 PID: 130 Comm: sh Not tainted 4.1.46 #2
[ 460.442026] Hardware name: Freescale i.MX6 Quad/DualLite (Device Tree)
[ 460.443113] Backtrace:
[ 460.443772] [<80013330>] (dump_backtrace) from [<80013544>] (show_stack+0x18/0x1c)
[...]
```

```
[ 460.476990] Unable to handle kernel NULL pointer dereference at virtual address 00000000
[ 460.477819] pgd = e71c8000
[ 460.478133] [00000000] *pgd=771ad831, *pte=00000000, *pppte=00000000
[ 460.479304] Internal error: Oops: 817 [#1] SMP ARM
[ 460.480044] Modules linked in:
[ 460.481368] CPU: 0 PID: 130 Comm: sh Not tainted 4.1.46 #2
[ 460.481945] Hardware name: Freescale i.MX6 Quad/DualLite (Device Tree)
[ 460.482577] task: e7abd000 ti: e728e000 task.ti: e728e000
[ 460.483090] PC is at sysrq_handle_crash+0x3c/0x4c
[ 460.483490] LR is at sysrq_handle_crash+0x34/0x4c
[ 460.483850] pc : [<8031bd80>]   lr : [<8031bd78>]   psr: a0000013
[ 460.483850] sp : e728fe50 ip : e728fe50 fp : e728fe5c
[ 460.484480] r10: 00000000 r9 : 00000000 r8 : 00000007
[ 460.484849] r7 : 00000000 r6 : 00000063 r5 : 80a4d5e8 r4 : 80a61694
[ 460.485287] r3 : 00000000 r2 : 00000001 r1 : f0004000 r0 : 00000063
[ 460.485741] Flags: NzCv IRQs on FIQs on Mode SVC_32 ISA ARM Segment user
[ 460.486446] Control: 10c5387d Table: 771c8059 DAC: 00000015
[ 460.486846] Process sh (pid: 130, stack limit = 0xe728e210)
[ 460.487240] Stack: (0xe728fe50 to 0xe7290000)
[ 460.487605] fe40:                                e728fe94 e728fe60 8031c744 8031bd50
[...]
```

```
[ 460.492794] Backtrace:
[ 460.493127] [<8031bd44>] (sysrq_handle_crash) from [<8031c744>] (__handle_sysrq+0xe0/0x254)
[ 460.493599] [<8031c664>] (__handle_sysrq) from [<8031cd14>] (write_sysrq_trigger+0x50/0x60)
[ 460.493982] r8:00000000 r7:e7bf5c00 r6:00000000 r5:00ab6400 r4:00000002
[ 460.494647] [<8031ccc4>] (write_sysrq_trigger) from [<8015815c>] (proc_reg_write+0x68/0x90)
```

```

[ 460.495126] r5:00000001 r4:00000000
[ 460.495537] [<801580f4>] (proc_reg_write) from [<800faa4c>] (__vfs_write+0x2c/0xe0)
[ 460.495945] r9:00ab6400 r8:00000002 r7:e728ff78 r6:e71aa8c0 r5:00ab6400 r4:80766b40
[ 460.496617] [<800faa20>] (__vfs_write) from [<800fb2f4>] (vfs_write+0x9c/0x15c)
[ 460.497292] r8:00000002 r7:00000002 r6:e728ff78 r5:00ab6400 r4:e71aa8c0
[ 460.498014] [<800fb258>] (vfs_write) from [<800fbb24>] (SyS_write+0x44/0x98)
[ 460.498402] r9:00ab6400 r8:00000002 r7:e71aa8c0 r6:e71aa8c0 r5:00000000 r4:00000000
[ 460.499154] [<800fbae0>] (SyS_write) from [<8000f960>] (ret_fast_syscall+0x0/0x54)
[ 460.499635] r9:e728e000 r8:8000fb44 r7:00000004 r6:00ab6400 r5:00000001 r4:000f8e2c
[ 460.500640] Code: 0a000000 e12fff33 e3a03000 e3a02001 (e5c32000)
[ 460.503826] Loading crashdump kernel...
[ 460.504424] Bye!
[ 0.000000] Booting Linux on physical CPU 0x0 << the dump kernel starts! >>
[ 0.000000] Linux version 4.1.46 (kai@klaptop) (gcc version 4.8.3 20140320 (prerelease)
(Sourcery CodeBench Lite 2014.05-29) ) #2 SMP Mon Nov 27 17:16:22 IST 2017
[ 0.000000] CPU: ARMv7 Processor [410fc090] revision 0 (ARMv7), cr=10c5387d
[ 0.000000] CPU: PIPT / VIPT nonaliasing data cache, VIPT nonaliasing instruction cache
[ 0.000000] Machine model: Freescale i.MX6 DualLite SABRE Smart Device Board
[ 0.000000] Ignoring memory block 0x10000000 - 0x50000000
[ 0.000000] cma: Reserved 16 MiB at 0x7ec00000
[ 0.000000] Memory policy: Data cache writeback
[ 0.000000] CPU: All CPU(s) started in SVC mode.
[ 0.000000] PERCPU: Embedded 12 pages/cpu @87cb9000 s16640 r8192 d24320 u49152
[ 0.000000] Built 1 zonelists in Zone order, mobility grouping on. Total pages: 32260
[ 0.000000] Kernel command line: console=ttyMXC0 root=/dev/mmcblk0 rootfstype=ext4 rootwait
init=/sbin/init maxcpus=1 reset_devices elfcorehdr=0x7ff00000 mem=130048K
[ 0.000000] PID hash table entries: 512 (order: -1, 2048 bytes)
...
[ 0.000000] Virtual kernel memory layout:
[ 0.000000] vector : 0xffff0000 - 0xffff1000 ( 4 kB)
[ 0.000000] fixmap : 0xffc00000 - 0xffff0000 (3072 kB)
[ 0.000000] vmalloc : 0x88000000 - 0xff000000 (1904 MB)
[ 0.000000] lowmem : 0x80000000 - 0x87f00000 ( 127 MB)
[ 0.000000] pkmap : 0x7fe00000 - 0x80000000 ( 2 MB)
[ 0.000000] modules : 0x7f000000 - 0x7fe00000 ( 14 MB)
[ 0.000000] .text : 0x80008000 - 0x809d7fdc (10048 kB)
[ 0.000000] .init : 0x809d8000 - 0x80a3a000 ( 392 kB)
[ 0.000000] .data : 0x80a3a000 - 0x80a9a6e0 ( 386 kB)
[ 0.000000] .bss : 0x80a9a6e0 - 0x812c3164 (8355 kB)
[ 0.000000] SLUB: HWalign=64, Order=0-3, MinObjects=0, CPUs=2, Nodes=1
[...]
```

```

[ 8.142812] fec 2188000.ethernet eth0: Link is Up - 100Mbps/Full - flow control rx/tx
[ 8.145506] IPv6: ADDRCONF(NETDEV_CHANGE): eth0: link becomes ready
[ 10.488149] cfg80211: Calling CRDA to update world regulatory domain

ARM / $ ls -lh /proc/vmcore
-r----- 1 0 0 1.9G Jan 1 00:07 /proc/vmcore
ARM / $ cp /proc/vmcore /kdump.img
ARM / $

```

(Very large dumpfile because the RAM on this system is 2 GB).

kdump Downsides

- A second kernel needs to be started when crashing
- Not all drivers work fine in the second kernel
- Very limited memory for the second kernel
- We need to construct a new initrd for the second kernel

Linux crash utility

Resources

Crash Whitepaper : superb!

https://crash-utility.github.io/crash_whitepaper.html

[Analyzing Linux kernel crash dumps with crash – Dedoimedo](#)

[‘crash’ source repo on GitHub](#) – do read the README here

[Using pstore to collect crash dumps – kernel-hardening mailing list, Kees Cook \[03Oct2019\]](#)

http://docs.oracle.com/cd/E37670_01/E41138/html/ch10s02.html

From the README

“ ...

At this point, x86, ia64, x86_64, ppc64, ppc, arm, arm64, alpha, mips, s390 and s390x-based kernels are supported.

...

- o One size fits all -- the utility can be run on any Linux kernel version dating back to 2.2.5-15. A primary design goal is to always maintain backwards-compatibility.
- o In order to contain debugging data, the top-level kernel Makefile's CFLAGS definition must contain the -g flag. Typically distributions will contain a package containing a vmlinux file with full debuginfo data. If not, the kernel must be rebuilt

...

The crash binary can only be used on systems of the same architecture as the host build system. There are a few optional manners of building the crash binary:

- o On an x86_64 host, a 32-bit x86 binary that can be used to analyze 32-bit x86 dumpfiles may be built by typing "make target=X86".
- o On an x86 or x86_64 host, a 32-bit x86 binary that can be used to analyze 32-bit arm dumpfiles may be built by typing "make target=ARM".
- o On an x86 or x86_64 host, a 32-bit x86 binary that can be used to analyze 32-bit mips dumpfiles may be built by typing "make target=MIPS".
- o On an ppc64 host, a 32-bit ppc binary that can be used to analyze

32-bit ppc dumpfiles may be built by typing "make target=PPC".

- o On an x86_64 host, an x86_64 binary that can be used to analyze arm64 dumpfiles may be built by typing "make target=ARM64".

...
”

Live System

Required: the vmlinux kernel image built with debug symbolic information.

Getting the kernel vmlinux with debug symbolic info

For Ubuntu

Where can one get the Ubuntu linux debug kernel vmlinux from? See:

https://wiki.ubuntu.com/Kernel/Systemtap#Where_to_get_debug_symbols_for_kernel_X.3F

Short answer: here:

<http://ddebs.ubuntu.com/pool/main/l/linux/>

Download the

linux-image-\$(uname -r)-dbgsym_\$(uname -r)_amd64.ddeb

file (assuming you're running on an x86_64 system).

Eg. Required file for kernel ver 3.16.0-37-generic is

linux-image-3.16.0-37-generic-dbgsym_3.16.0-37.49_amd64.ddeb

Extract the vmlinux-\$(uname -r) image from the downloaded ddeb file.

Unresolved: recent Ubuntu x86_64 (tried with 16.10, 17.04 and 17.10), there seems to be an issue running crash with the Ubuntu debugsym vmlinux (it *does* work well with older Ubuntu distros; tried with 14.04 LTS and it's fine). ?? **<< Update: does work on Ubuntu 18.04.2 LTS >>**

For Fedora

Kernel-debug repo:

<https://www.rpmfind.net/linux/rpm2html/search.php?query=kernel-debug>

How to use kdump to debug kernel crashes

https://fedoraproject.org/wiki/How_to_use_kdump_to_debug_kernel_crashes

Essentially, need to do this to get the latest debug kernel:

`dnf install kernel-debuginfo`

(Observation: it's not always *the* latest kernel installed, so debugging the “live” way becomes an issue).

“Live” system investigation with crash

sudo crash <vmlinux-with-symbolic-info> <kdump-image> [corr System.map]

Live debug

sudo crash <vmlinux-with-symbolic-info> /proc/kcore [corr System.map]

For a live debug session, the `<vmlinux-with-symbolic-info>` must precisely match the currently running kernel version

```
$ sudo crash vmlinux-3.16.0-37-generic /proc/kcore /boot/System.map-3.16.0-37-generic
```

crash 7.0.8

Copyright (C) 2002-2014 Red Hat, Inc.

Copyright (C) 2004, 2005, 2006, 2010 IBM Corporation

Copyright (C) 1999-2006 Hewlett-Packard Co

Copyright (C) 2005, 2006, 2011, 2012 Fujitsu Limited

Copyright (C) 2006, 2007 VA Linux Systems Japan K.K.

Copyright (C) 2005, 2011 NEC Corporation

Copyright (C) 1999, 2002, 2007 Silicon Graphics, Inc.

Copyright (C) 1999, 2000, 2001, 2002 Mission Critical Linux, Inc.

This program is free software, covered by the GNU General Public License,

and you are welcome to change it and/or distribute copies of it under

certain conditions. Enter "help copying" to see the conditions.

This program has absolutely no warranty. Enter "help warranty" for details.

GNU gdb (GDB) 7.6

Copyright (C) 2013 Free Software Foundation, Inc.

License GPLv3+: GNU GPL version 3 or later <<http://gnu.org/licenses/gpl.html>>

This is free software: you are free to change and redistribute it.

There is NO WARRANTY, to the extent permitted by law. Type "show copying"

and "show warranty" for details.

This GDB was configured as "x86_64-unknown-linux-gnu"...

SYSTEM MAP: /boot/System.map-3.16.0-37-generic

DEBUG KERNEL: /home/kaiwan/Downloads/vmlinux-3.16.0-37-generic (3.16.0-37-generic)

DUMPFILE: /proc/kcore

CPUS: 4

DATE: Wed May 13 21:50:24 2015

UPTIME: 2 days, 01:47:47

LOAD AVERAGE: 2.26, 2.40, 2.76

TASKS: 710

NODENAME: kaiwan-ThinkPad-X220

RELEASE: 3.16.0-37-generic

VERSION: #51-Ubuntu SMP Tue May 5 13:45:59 UTC 2015

MACHINE: x86_64 (2691 Mhz)

MEMORY: 7.9 GB

PID: 25671

COMMAND: "crash"

TASK: fffff80019cf0000 [THREAD_INFO: fffff8000b1fc000]

CPU: 3

STATE: TASK_RUNNING (ACTIVE)

crash>

Crash Commands - Quick Notes

The tool's environment is context-specific. On a live system, the default context is the command itself; on a dump the default context will be the task that panicked [can be changed with the 'set' command].

Structures

Do what	How / Command	Example
See structure definition	Give name of structure or use ' whatis ' <struct_name>	crash> file crash> whatis file

- disassemble
 - disassemble /r : print opcode and instruction
 - rd : read memory from address
 - gdb list *(memcpy+16): find c-code line
 - mod: print modules
 - no information for module loaded: ffffffff0316fa0 kvm 342174 (not loaded) [CONFIG_KALLSYMS]
 - mod -S: find modules /lib/modules//
 - There not only vmlinux but also modules should be copied
-

Common Commands

<u>bt</u>	<p>Display the backtrace of the current context, or as specified with arguments. This command is typically the first command entered after starting a dumpfile session. Since the initial context is the panic context, it will show the function trace leading up to the kernel panic. bt -a will show the trace of the <i>active</i> task on each CPU, since there may be an interrelationship between the panicking task on one CPU and the running task(s) on the other CPU(s). When bt is given as the argument to <u>foreach</u>, displays the backtraces of <i>all</i> tasks.</p> <p><< <i>bt -l</i> : show file and line number of each stack trace text location. >></p>
<u>struct</u>	<p>Print the contents of a data structure at a specified address. This command is so common that it is typically unnecessary to enter the struct command name on the command line; if the first command line argument is not a crash or gdb command, but it is the name of a known data structure, then all the command line arguments are passed to the struct command. So for example, the following two commands yield the same result:</p> <pre>crash> struct vm_area_struct d3cb2600</pre> <pre>crash> vm_area_struct d3cb2600</pre>
<u>set</u>	<p>Set a new task context by PID, task address, or cpu. Since several crash commands are context-sensitive, it's helpful to be able to change the context to avoid having to pass the PID or task address to those context-sensitive commands in order to access the data of a task that is <i>not</i> the current context.</p>
<u>p</u>	<p>Prints the contents of a kernel variable; since it's a gateway to the print command of the mbedded gdb module, it can also be used to print complex C language expressions.</p>
<u>rd</u>	<p>Read memory, which may be either kernel virtual, user virtual, or</p>

	physical, and display it several different formats and sizes.
<u>ps</u>	Lists basic task information for each process; it can also display parent and child hierarchies.
<u>log</u>	Dump the kernel <code>log_buf</code> , which often contains clues leading up to a subsequent kernel crash.
<u>foreach</u>	Execute a <code>crash</code> command on all tasks, or those specified, in the system; can be used with <u>bt</u> , <u>vm</u> , <u>task</u> , <u>files</u> , <u>net</u> , <u>set</u> , <u>sig</u> and <u>vtop</u> .
<u>files</u>	Dump the open file descriptor data of a task; most usefully, the <code>file</code> , <code>dentry</code> and <code>inode</code> structure addresses for each open file descriptor.
<u>vm</u>	Dump the virtual memory map of a task, including the vital information concerning each <code>vm_area_struct</code> making up a task's address space. It can also dump the physical address of each page in the address space, or if not mapped, its location in a file or on the swap device.

whatis <struct-or-symbol-name> : displays the structure members

Eg.

<< on Ubuntu 15.04 kernel ver 3.19.0-22-generic >>

```
crash> whatis thread_info
struct thread_info {
    struct task_struct *task;
    struct exec_domain *exec_domain;
    __u32 flags;
    __u32 status;
    __u32 cpu;
    int saved_preempt_count;
    mm_segment_t addr_limit;
    struct restart_block restart_block;
    void *sysenter_return;
    unsigned int sig_on_uaccess_error : 1;
    unsigned int uaccess_err : 1;
}
SIZE: 104
crash> struct task_struct {
    volatile long state;
    void *stack;
    atomic_t usage;
    ...
    unsigned int sequential_io_avg;
}
SIZE: 2504
crash>
```

Disassembly:

```
crash> dis printk
0xfffffffff817c1bed <printk>:    data32 data32 data32 xchg %ax,%ax [FTRACE NOP]
0xfffffffff817c1bf2 <printk+5>:  push   %rbp
0xfffffffff817c1bf3 <printk+6>:  mov    %rsp,%rbp
0xfffffffff817c1bf6 <printk+9>:  sub    $0x50,%rsp
...
```

Disassemble 20 instructions of `tcp_sendmsg`, showing source line numbers (`-l` option) and change radix to hex (`-x`):

```
crash> dis -l tcp_sendmsg 20 -x
/build/buildd/linux-3.19.0/net/ipv4/tcp.c: 1069
0xfffffffff8170bf80 <tcp_sendmsg>:    data32 data32 data32 xchg %ax,%ax [FTRACE NOP]
0xfffffffff8170bf85 <tcp_sendmsg+0x5>:  push   %rbp
0xfffffffff8170bf86 <tcp_sendmsg+0x6>:  mov    %rsp,%rbp
0xfffffffff8170bf89 <tcp_sendmsg+0x9>:  push   %r15
0xfffffffff8170bf8b <tcp_sendmsg+0xb>:  push   %r14
0xfffffffff8170bf8d <tcp_sendmsg+0xd>:  push   %r13
0xfffffffff8170bf8f <tcp_sendmsg+0xf>:  push   %r12
0xfffffffff8170bf91 <tcp_sendmsg+0x11>: mov    %rsi,%r15
0xfffffffff8170bf94 <tcp_sendmsg+0x14>: push   %rbx
/build/buildd/linux-3.19.0/include/net/sock.h: 1536
0xfffffffff8170bf95 <tcp_sendmsg+0x15>: xor    %esi,%esi
/build/buildd/linux-3.19.0/net/ipv4/tcp.c: 1069
0xfffffffff8170bf97 <tcp_sendmsg+0x17>: mov    %rdx,%rbx
/build/buildd/linux-3.19.0/include/net/sock.h: 1536
0xfffffffff8170bf9a <tcp_sendmsg+0x1a>: mov    %r15,%rdi
/build/buildd/linux-3.19.0/net/ipv4/tcp.c: 1069
0xfffffffff8170bf9d <tcp_sendmsg+0x1d>: mov    %rcx,%r12
0xfffffffff8170bfa0 <tcp_sendmsg+0x20>: sub    $0x98,%rsp
...
```

Running crash in “batch mode”

Very useful technique to grab information from a dumpfile and save it.

1. Create a crash “commands script” file:

```
$ cat crash_getinfo
echo "=== System Info ==="
echo "--- sys ---"
sys

echo "--- log ---"
log

echo "--- ps ---"
ps

echo "--- dev ---"
dev

echo "=== Current Context Info ==="
echo "--- bt -a ---"
bt -a

echo "--- files ---"
files
```

```
echo "--- vm ---"
vm

exit
$
```

2. Invoke it via crash:

```
sudo crash dumpfile.img vmlinux_dbgsym.img < crash_getinfo > report.txt
```

sym (symbol)

[Running crash on a dump image from an ARM-32; [see how to here](#)]

```
...
crash_32bit_for_arm> help sym
NAME
  sym - translate a symbol to its virtual address, or vice-versa
```

SYNOPSIS

```
sym [-l] | [-M] | [-m module] | [-p|-n] | [-q string] | [symbol | vaddr]
```

DESCRIPTION

This command translates a symbol to its virtual address, or a static kernel virtual address to its symbol -- or to a symbol-plus-offset value, if appropriate. Additionally, the symbol type is shown in parentheses, and if the symbol is a known text value, the file and line number are shown.

```
...
<lots of examples!>
```

```
crash_32bit_for_arm> bt
```

```
PID: 735 TASK: 9f6af900 CPU: 0 COMMAND: "echo"
#0 [<804060d8>] (sysrq_handle_crash) from [<804065bc>]
#1 [<804065bc>] (__handle_sysrq) from [<80406ab8>]
#2 [<80406ab8>] (write_sysrq_trigger) from [<80278588>]
#3 [<80278588>] (proc_reg_write) from [<802235c4>]
#4 [<802235c4>] (__vfs_write) from [<80224098>]
#5 [<80224098>] (vfs_write) from [<80224d30>]
#6 [<80224d30>] (sys_write) from [<801074a0>]
pc : [<76e8d7ec>] lr : [<0000f9dc>] psr: 60000010
sp : 7ebdcc7c ip : 00000000 fp : 00000000
r10: 0010286c r9 : 7ebdce68 r8 : 00000020
r7 : 00000004 r6 : 00103008 r5 : 00000001 r4 : 00102e2c
r3 : 00000000 r2 : 00000002 r1 : 00103008 r0 : 00000001
Flags: nZCv IRQs on FIQs on Mode USER_32 ISA ARM
```

```
crash_32bit_for_arm> sym 801074a0
```

```
801074a0 (t) ret_fast_syscall ../arch/arm/kernel/entry-common.S
```

```
crash_32bit_for_arm> bt -l << -l: show file and line number info >>
```

```
PID: 735 TASK: 9f6af900 CPU: 0 COMMAND: "echo"
#0 [<804060d8>] (sysrq_handle_crash) from [<804065bc>]
<...>/linux-4.9.1/drivers/tty/sysrq.c: 144
#1 [<804065bc>] (__handle_sysrq) from [<80406ab8>]
<...>/linux-4.9.1/drivers/tty/sysrq.c: 552
#2 [<80406ab8>] (write_sysrq_trigger) from [<80278588>]
<...>/linux-4.9.1/drivers/tty/sysrq.c: 1101
#3 [<80278588>] (proc_reg_write) from [<802235c4>]
<...>/linux-4.9.1/fs/proc/inode.c: 216
#4 [<802235c4>] (__vfs_write) from [<80224098>]
<...>/linux-4.9.1/fs/read_write.c: 510
#5 [<80224098>] (vfs_write) from [<80224d30>]
```

```

<...>/linux-4.9.1/fs/read_write.c: 561
#6 [<80224d30>] (sys_write) from [<801074a0>]
<...>/linux-4.9.1/fs/read_write.c: 608
pc : [<76e8d7ec>]   lr : [<0000f9dc>]   psr: 60000010
sp : 7ebdcc7c   ip : 00000000   fp : 00000000
r10: 0010286c   r9 : 7ebdce68   r8 : 00000020
r7 : 00000004   r6 : 00103008   r5 : 00000001   r4 : 00102e2c
r3 : 00000000   r2 : 00000002   r1 : 00103008   r0 : 00000001
Flags: nZCv   IRQs on   FIQs on   Mode USER_32   ISA ARM
crash_32bit_for_arm>

```

Module Debugging

Compile the kernel module with -g (edit it's Makefile, specify
EXTRA_CFLAGS += -DDEBUG -g

mod -S :
Load the symbolic and debugging data of all modules

Appending the directory pathname (where your kernel module resides) will restrict the search to that folder...

Eg.

```

crash> mod -S /home/seawolf/kaiwanTECH/L3_dd_trg/miscmj_tst/
MODULE      NAME      SIZE  OBJECT FILE
ffffffffc0393480  pata_acpi      16384 /lib/modules/4.17.0/kernel/drivers/ata/pata_acpi.ko
ffffffffc039d500  i2c_piix4      24576 /lib/modules/4.17.0/kernel/drivers/i2c/busses/i2c-piix4.ko
ffffffffc03a9580  libahci        32768 /lib/modules/4.17.0/kernel/drivers/ata/libahci.ko

[...]

ffffffffc072a940  vboxsf         45056 /lib/modules/4.17.0/misc/vboxsf.ko
ffffffffc075c080  miscmj_tst     16384 /home/seawolf/kaiwanTECH/L3_dd_trg/miscmj_tst/miscmj_tst.o
crash>

```

Set ctx to the process accessing the driver:

```

crash> set 4675
PID: 4675
COMMAND: "echo"
TASK: ffff9e7dbb865a00 [THREAD_INFO: ffff9e7dbb865a00]
CPU: 1
STATE: TASK_INTERRUPTIBLE

```

Now dump the stack frames - all show up!!

```

crash> bt
PID: 4675 TASK: ffff9e7dbb865a00 CPU: 1 COMMAND: "echo"
#0 [ffffc24c01607cd0] __schedule at fffffffffa133731b
#1 [ffffc24c01607d68] schedule at fffffffffa13378dc
#2 [ffffc24c01607d78] schedule_timeout at fffffffffa133b69b
#3 [ffffc24c01607df8] my_dev_write at fffffffffc075a0ff [miscmj_tst]
#4 [ffffc24c01607e18] __vfs_write at fffffffffa0c8bd3a
#5 [ffffc24c01607ea0] vfs_write at fffffffffa0c8c011

```

```
#6 [ffffc24c01607ed8] ksys_write at ffffffffafa0c8c2a5
#7 [ffffc24c01607f20] __x64_sys_write at ffffffffafa0c8c32a
#8 [ffffc24c01607f30] do_syscall_64 at ffffffffafa0a041fa
#9 [ffffc24c01607f50] entry_SYSCALL_64_after_hwframe at ffffffffafa1400088
...
```

crash>

bt -f

```
...
#3 [ffffc24c01607df8] my_dev_write at ffffffffcc075a0ff [miscmj_tst]
ffffc24c01607e00: 0000000000000000 ffff9e7db6d27900
ffffc24c01607e10: fffffc24c01607e98 ffffffffafa0c8bd3a
#4 [ffffc24c01607e18] __vfs_write at ffffffffafa0c8bd3a
ffffc24c01607e20: ffff9e7db88ad330 fffff0cdc0e22b70
ffffc24c01607e30: 0000000000000000 0dabdbe391c9e100
ffffc24c01607e40: 0000000000000055 ffff9e7db46c2708
ffffc24c01607e50: 00005643ea666408 ffff9e7db6e49080
ffffc24c01607e60: 0000000000000040 fffffc24c01607ea0
ffffc24c01607e70: ffffffffafa0c21fc3 0000000000000006
ffffc24c01607e80: 0dabdbe391c9e100 0000000000000004
ffffc24c01607e90: 0000000000000000 fffffc24c01607ed0
ffffc24c01607ea0: ffffffffafa0c8c011
#5 [ffffc24c01607ea0] vfs_write at ffffffffafa0c8c011
ffffc24c01607ea8: ffff9e7db6d27900 ffff9e7db6d27900
ffffc24c01607eb8: 00005643ea665400 0000000000000004
ffffc24c01607ec8: 0000000000000000 fffffc24c01607f18
ffffc24c01607ed8: ffffffffafa0c8c2a5
...
```

Note-

- [SETTING UP KDUMP AND CRASH FOR ARM-32 – AN ONGOING SAGA, kaiwanTECH, July 2017](#)
- [Running crash on ARM- see this link from the Linaro Wiki](#)
- **crash** used quite a bit here: [A Short Guide to Kernel Debugging: A story about finding a kernel bug on a production system, Square](#)
- **Unresolved:** recent Ubuntu x86_64 (tried with 16.10, 17.04 and 17.10), there seems to be an issue running crash with the Ubuntu debugsym vmlinux (it *does* work well with older Ubuntu distros; tried with 14.04 LTS and it's fine). ??
<< Update: does work on Ubuntu 18.04.2 LTS >>

<< End document >>